

Generalized Flow (gflow) as a Loss Function

Luke Hassel

February 23, 2026

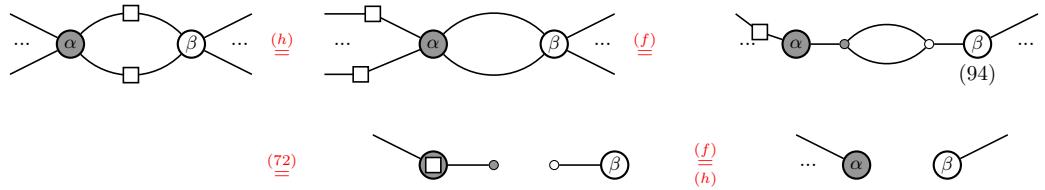
0.1 Graph-like ZX-Diagrams

A ZX-diagram is called *graph-like* [3] if:

- All spiders are Z-spiders (green).
- All edges are Hadamard edges (dashed blue lines).
- There are no self-loops or parallel edges between the same pair of spiders.
- Every input or output is connected to a distinct spider, and no spider is connected to more than one input or output.

The conversion from a ZX-diagram to a graph like ZX-diagrams is done as follows [3]:

1. Convert all X-spiders into Z-spiders by pushing out Hadamards (by using **(h)** of Figure 1).
2. Cancel all adjacent pairs of Hadamard gates using **(hh)**.
3. Fuse all connected spiders using **(f)**.
4. View the remaining Hadamard gates between spiders as the Hadamard-edges of **(92)**.
5. Remove all self-loops using **(52)** and **(83)**.
6. If two spiders are connected by multiple Hadamard-edges, a variation of the Hopf-rule removes these:



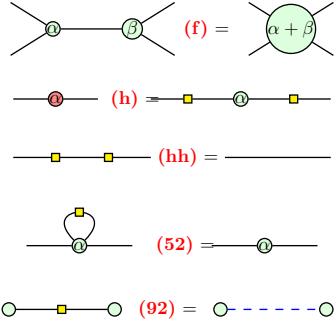


Figure 1: Key ZX-calculus rules used in the transformation: Fusion (f), Color Change (h), Hadamard Cancellation (hh), Loop Removal (52), and Hadamard Edge Definition (92).

After the conversion the ZX-diagram only contains Z-spiders and Hadamard-edges.

0.2 Open Graphs

An *open graph* (G, I, O, λ) is an undirected graph $G = (V, E)$ where $I \subseteq V$ are the inputs and $O \subseteq V$ are the outputs. [1].

A ZX-diagram can be converted to an open graph. Each Z-spider in a graph like ZX-diagram corresponds to a vertex $v \in V$. The input set I consists of all vertices connected to an input wire, and the output set O consists of all vertices connected to an output wire. The edges E of the open graph correspond exactly to the Hadamard edges between the spiders in the diagram [1].

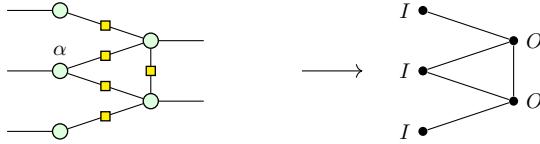


Figure 2: Conversion from a graph-like ZX-diagram (left) to an open graph (right). Spiders with input wires become the input set I , and spiders with output wires become the output set O . Hadamard edges become graph edges.

0.3 Focused Generalized Flow

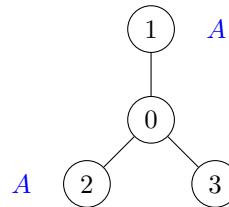
The authors of [1] introduce a special case of the generalised flow definition used in MBQC called focused generalised flow since as a graph property the two notations are equivalent.

(Focused gFlow): Given an open graph $G = (V, E, I, O)$, a *focused gFlow* (g, \prec) on G consists of a function $g : \bar{O} \rightarrow 2^I$ and a partial order \prec on the vertices V of G such that for all $u \in \bar{O}$:

1. $\text{Odd}_G(g(u)) \cap \bar{O} = \{u\}$
 2. $\forall v \in q(u), u \prec v$

where $\bar{O} = V \setminus O$ is the set of measured vertices, $\bar{I} = V \setminus I$ is the set of vertices available for the flow map, $2^{\bar{I}}$ is the powerset of \bar{I} , and $\text{Odd}_G(A) := \{v \in V(G) \mid |N(v) \cap A| \equiv 1 \pmod{2}\}$ is the odd neighbourhood of A .

Example: Calculating Odd_G(A) Consider the star graph G with center vertex 0 and leaves $\{1, 2, 3\}$. Let the subset $A = \{1, 2\}$.



To find $\text{Odd}_G(\{1, 2\})$, check the parity of neighbors for each vertex:

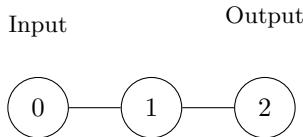
- **Vertex 0:** Neighbors are $\{1, 2, 3\}$. Intersection with A is $\{1, 2\}$. Count = 2 (even). $\rightarrow 0 \notin \text{Odd}_G(A)$.
 - **Vertex 1:** Neighbor is $\{0\}$. Intersection with A is \emptyset . Count = 0 (even). $\rightarrow 1 \notin \text{Odd}_G(A)$.
 - **Vertex 2:** Neighbor is $\{0\}$. Intersection with A is \emptyset . Count = 0 (even). $\rightarrow 2 \notin \text{Odd}_G(A)$.
 - **Vertex 3:** Neighbor is $\{0\}$. Intersection with A is \emptyset . Count = 0 (even). $\rightarrow 3 \notin \text{Odd}_G(A)$.

In this case, $\text{Odd}_G(\{1, 2\}) = \emptyset$. However, if $A = \{0\}$, then all neighbors $\{1, 2, 3\}$ have exactly one neighbor in A , so $\text{Odd}_G(\{0\}) = \{1, 2, 3\}$.

1 GFlow Examples

1.1 Example 1: Linear Chain (Has GFlow)

Consider a 3-qubit linear chain with $I = \{0\}$ and $O = \{2\}$:



GFlow:

$$g(0) = \{1\} \quad g(1) = \{2\}$$

Partial order: $0 \prec 1 \prec 2$.

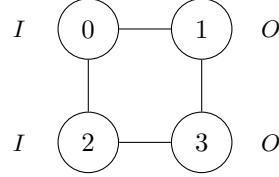
Verification of Condition 1 ($\text{Odd}_G(g(u)) \cap \bar{O} = \{u\}$): Here, the non output set is $\bar{O} = V \setminus O = \{0, 1\}$.

- **For $u = 0$:** $g(0) = \{1\}$. The neighbors of 1 are $\{0, 2\}$, so $\text{Odd}_G(g(0)) = \{0, 2\}$.
Intersection: $\{0, 2\} \cap \{0, 1\} = \{0\}$. ✓
- **For $u = 1$:** $g(1) = \{2\}$. The neighbor of 2 is $\{1\}$, so $\text{Odd}_G(g(1)) = \{1\}$.
Intersection: $\{1\} \cap \{0, 1\} = \{1\}$. ✓

Condition 2 ($\forall v \in g(u), u \prec v$) is satisfied as $0 \prec 1$ and $1 \prec 2$.

1.2 Example 2: 2×2 Grid (Has GFlow)

Consider a 4-qubit grid (box) with two inputs $I = \{0, 2\}$ and two outputs $O = \{1, 3\}$:



GFlow:

$$g(0) = \{1\} \quad g(2) = \{3\}$$

Partial order: $0 \prec 1$ and $2 \prec 3$.

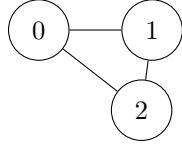
Verification of Condition 1 ($\text{Odd}_G(g(u)) \cap \bar{O} = \{u\}$): The measured set is $\bar{O} = V \setminus O = \{0, 2\}$.

- **For $u = 0$:** $g(0) = \{1\}$. The neighbors of 1 are $\{0, 3\}$, so $\text{Odd}_G(g(0)) = \{0, 3\}$.
Intersection: $\{0, 3\} \cap \{0, 2\} = \{0\}$. ✓
- **For $u = 2$:** $g(2) = \{3\}$. The neighbors of 3 are $\{1, 2\}$, so $\text{Odd}_G(g(2)) = \{1, 2\}$.
Intersection: $\{1, 2\} \cap \{0, 2\} = \{2\}$. ✓

1.3 Example 3: Triangle (No GFlow)

Consider a triangle with $I = \{0\}$, $O = \{0\}$ (input and output are the same vertex):

I/O



Any assignment creates a cycle: $g(1)$ must use $\{2\}$ and $g(2)$ must use $\{1\}$, requiring both $1 \prec 2$ and $2 \prec 1$, which violates the partial order. **No gflow exists.**

2 Mhalla-Perdrix Algorithm

The standard algorithm for finding a gflow is the algorithm at [2], which operates in $O(n^3)$ time. It relies on solving linear systems over the field \mathbb{F}_2 .

The core logic involves:

1. Constructing the adjacency matrix of the graph state.
2. Identifying candidate sets using Gaussian elimination over \mathbb{F}_2 .
3. Iteratively building the gflow layers from outputs back to inputs.

But the algorithm only works with values fixed to 0 or 1 in the g-assigments and because of its iterative nature it is not possible to use it to differentiate.

3 Differentiable Gflow Relaxation

3.1 Loss Function Construction

Since the graph is undirected, only the upper triangular part of the adjacency matrix is needed. Assuming an optimal assignment of g exists, let G^* denote its matrix form.

For example, considering the 2×2 grid from Example 2 with measured qubits $\bar{O} = \{0, 2\}$ and candidate qubits $\bar{I} = \{1, 3\}$, the optimal assignment $g(0) = \{1\}, g(2) = \{3\}$ corresponds to the flow matrix:

$$G^* = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \quad (1)$$

where the rows correspond to $u \in \{0, 2\}$ and columns to $v \in \{1, 3\}$.

3.1.1 Odd Loss

For a vertex u to be in \bar{O} , the g-set $g(u)$ must have an odd number of connections to u . This can be expressed as an XOR sum over all vertices v :

$$\bigoplus_{v \in V} A_{uv} G_{uv} = 1 \quad (2)$$

where $G_{uv} = 1$ if $v \in g(u)$ and 0 otherwise.

Example Parity Check Consider a vertex u with two physical neighbors v_1 and v_2 (so $A_{uv_1} = 1$, $A_{uv_2} = 1$, and $A_{uv} = 0$ for others). The XOR sum calculates the parity of the selected g-set vertices connected to u :

- If we select $g(u) = \{v_1\}$, then $G_{uv_1} = 1$ and $G_{uv_2} = 0$. The sum is $1 \oplus 0 = 1$. (Condition satisfied, u has an odd number of neighbors in the g-set).
- If we select $g(u) = \{v_1, v_2\}$, then $G_{uv_1} = 1$ and $G_{uv_2} = 1$. The sum is $1 \oplus 1 = 0$. (Condition failed, u has an even number of neighbors in the g-set).

This illustrates how G acts as a filter on the neighborhood defined by A .

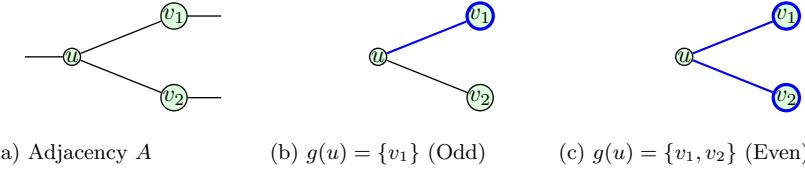


Figure 3: Open graph fragment illustrating the parity XOR sum. Blue thick edges indicate active neighbors in the g-set ($G_{uv} = 1$).

Note that this captures only the relaxed condition $u \in \text{Odd}_G(g(u))$. The full Focused GFlow definition also requires $\text{Odd}_G(g(u)) \cap \bar{O} = \{u\}$, which is not explicitly enforced.

Example: Focused Condition Violation Consider two non-output vertices $u, u' \in \bar{O}$ that share a common neighbor v , as shown in Figure 4. If the optimization chooses $g(u) = \{v\}$, the condition $u \in \text{Odd}_G(g(u))$ is satisfied since u has exactly one neighbor in its correction set. However, u' also has an odd number of neighbors in the same set $g(u)$, leading to $\text{Odd}_G(g(u)) \cap \bar{O} = \{u, u'\}$. The current loss L_{Odd} does not penalize this because it only evaluates the neighborhood of $g(u)$ at vertex u , failing to enforce that u must be the *unique* non-output vertex in that neighborhood.

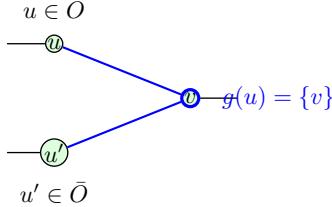


Figure 4: Violation of the focused condition. Selecting $g(u) = \{v\}$ satisfies the parity condition for u , but inadvertently creates an odd neighborhood for u' as well. Since both $u, u' \in \bar{O}$, the condition $\text{Odd}_G(g(u)) \cap \bar{O} = \{u\}$ is violated.

To also allow continuous values (including learning the graph structure $A_{uv} \in [0, 1]$), we map the XOR sum to a product over all vertices:

$$\prod_{v \in V} (1 - 2A_{uv}G_{uv}) = -1 \quad (3)$$

In the binary case, this product is -1 when the number of active neighbors (where $A_{uv}G_{uv} = 1$) is odd, and 1 when even. For natural continuous relaxation, the term $1 - 2A_{uv}G_{uv}$ interpolates smoothly: if an edge is absent ($A_{uv} \approx 0$), the factor is close to 1 and does not affect the product. Conversely, if an edge is present ($A_{uv} \approx 1$) but the flow weight is undecided ($G_{uv} \approx 0.5$), the factor vanishes.

Based on this product formulation, the parity loss is defined as the normalized squared distance from the target value -1 :

$$\frac{1}{4} \left(\prod_{v \in V} (1 - 2A_{uv}G_{uv}) + 1 \right)^2 \quad (4)$$

The term is squared to ensure non-negativity. The factor $1/4$ normalizes the loss: if the condition is satisfied (product is -1), the loss is $\frac{1}{4}(-1 + 1)^2 = 0$. If violated (product is 1), the loss is $\frac{1}{4}(1 + 1)^2 = 1$.

Expanding this to all nodes $u \in \bar{O}$ yields:

$$L_{\text{Odd}} = \frac{1}{4} \sum_{u \in \bar{O}} \left(\prod_{v \in V} (1 - 2A_{uv}G_{uv}) + 1 \right)^2 \quad (5)$$

3.1.2 Order Loss

To enforce the partial order condition $\forall v \in g(u), u \prec v$ a real-valued ‘‘order’’ value $\tau_v \in \mathbb{R}$ is assigned to each vertex. The discrete condition $u \prec v$ maps to the continuous constraint $\tau_u < \tau_v$.

It is important to note that the values τ are separate variables from the Flow map G and are optimized jointly to satisfy the causal constraints.

This constraint implies that for any active dependency where $G_{uv} \approx 1$ (indicating $v \in g(u)$), the time coordinate of u must be strictly less than that of v . Violations are penalized using a Rectified Linear Unit (ReLU) function with a margin $\epsilon > 0$:

$$L_{\text{order}} = \sum_{u,v} G_{uv} \cdot \text{ReLU}(\tau_u - \tau_v + \epsilon)^2 \quad (6)$$

Here, $\tau_u - \tau_v + \epsilon$ measures the violation. If $\tau_v > \tau_u + \epsilon$ (valid ordering), the term is negative and the ReLU equates to zero. If the ordering is violated or the separation is insufficient, the term is positive, creating a loss proportional to the strength of the assignment G_{uv} .

The margin ϵ is crucial because a standard $\text{ReLU}(\tau_u - \tau_v)$ would only enforce the non-strict inequality $\tau_u \leq \tau_v$. By adding ϵ , we ensure a minimum separation $\tau_v \geq \tau_u + \epsilon$, thereby guaranteeing the strict inequality $\tau_u < \tau_v$ required for a valid order.

The term is squared to penalize larger violations more severely.

Example Calculation Consider two vertices u and v where $G_{uv} = 1$ (meaning $v \in g(u)$), and let the margin be $\epsilon = 0.1$.

- **Valid Ordering:** If $\tau_u = 0.5$ and $\tau_v = 1.0$, then:

$$\text{ReLU}(0.5 - 1.0 + 0.1) = \text{ReLU}(-0.4) = 0$$

The loss contribution is zero, as the partial order $u \prec v$ is satisfied with sufficient separation.

- **Invalid Ordering:** If $\tau_u = 0.8$ and $\tau_v = 0.7$ (a backward dependency), then:

$$\text{ReLU}(0.8 - 0.7 + 0.1) = \text{ReLU}(0.2) = 0.2$$

The loss contribution is $(0.2)^2 = 0.04$, penalizing the violation.

3.2 Differentiability

For a given graph structure defined by the adjacency matrix A , we seek to identify a valid gflow G and ordering τ by minimizing the total loss:

$$L(A, G, \tau) = L_{\text{Odd}}(A, G) + L_{\text{Order}}(G, \tau) \quad (7)$$

The optimal gflow assignment G^* (and corresponding τ^*) is obtained by optimizing this objective:

$$G^*(A) = \arg \min_{G, \tau} L(A, G, \tau) \quad (8)$$

We then define the function $f(A)$ as the minimum loss value achieved:

$$f(A) = L(A, G^*(A), \tau^*(A)) \quad (9)$$

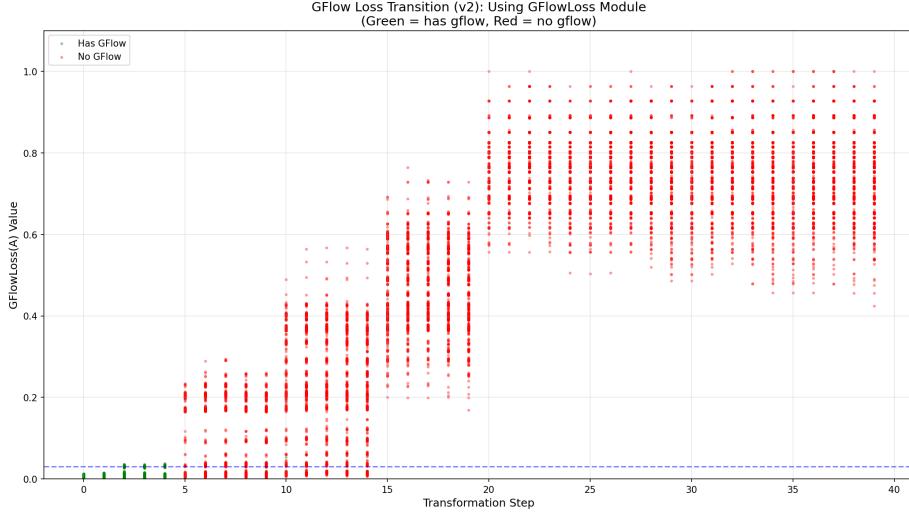


Figure 5: Loss with 6 nodes. Here local complementation is used to create different initial graphs.

Since the loss components are constructed from differentiable operations (element-wise products, summations, and ReLU functions), the resulting function $f(A)$ remains differentiable with respect to the input adjacency matrix A . This allows gradients to be backpropagated through the gflow optimization process to update the underlying graph generator efficiently.

3.3 Simple Gflow Model

Training a simple generative model takes a random vector and outputs an open graph based on the described loss function. This yields gflow in 94% of the open graphs. BFS transformation of the remaining graphs ensures 100% gflow existence.

3.4 Local complementation

Local complementation is used to create different initial graphs for the scatter plots shown in Figure 8. This operation preserves gflow existence [1]. Applying local complementation to a vertex v inverts the adjacency of its neighborhood $N(v)$. For any two neighbors $u, w \in N(v)$, the operation removes the edge (u, w) if it exists, or adds it if it is missing. This transformation leaves the adjacency of v and any vertex outside $N(v)$ unchanged.

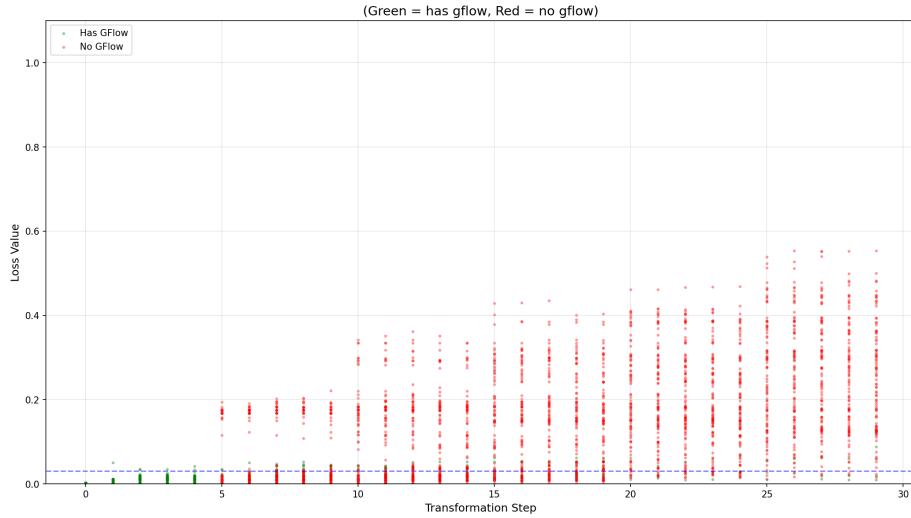


Figure 6: Loss with 20 nodes.

References

- [1] Ross Duncan, Aleks Kissinger, Simon Perdrix, and John van de Wetering. Graph-theoretic simplification of quantum circuits with the zx-calculus. *Quantum*, 4:279, 6 2020.
- [2] Mehdi Mhalla and Simon Perdrix. Finding optimal flows efficiently. *arXiv preprint arXiv:0709.2670*, 2008.
- [3] John van de Wetering. Zx-calculus for the working quantum computer scientist. *arXiv preprint arXiv:2012.13966*, 2020.

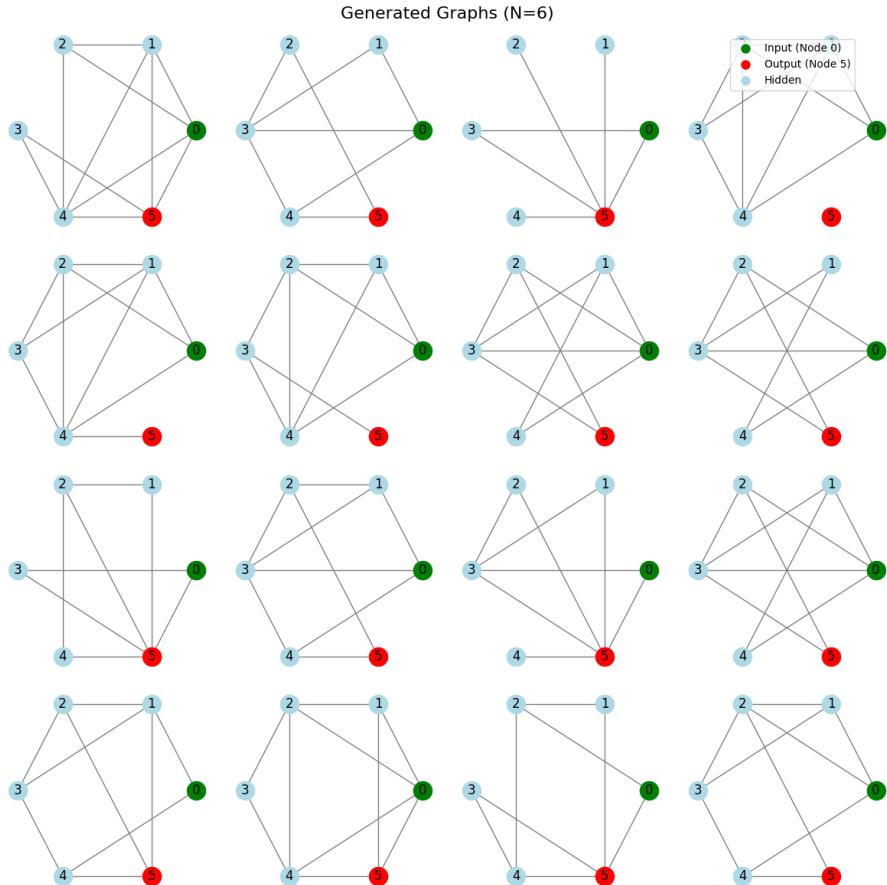


Figure 7: Generator output visualization.

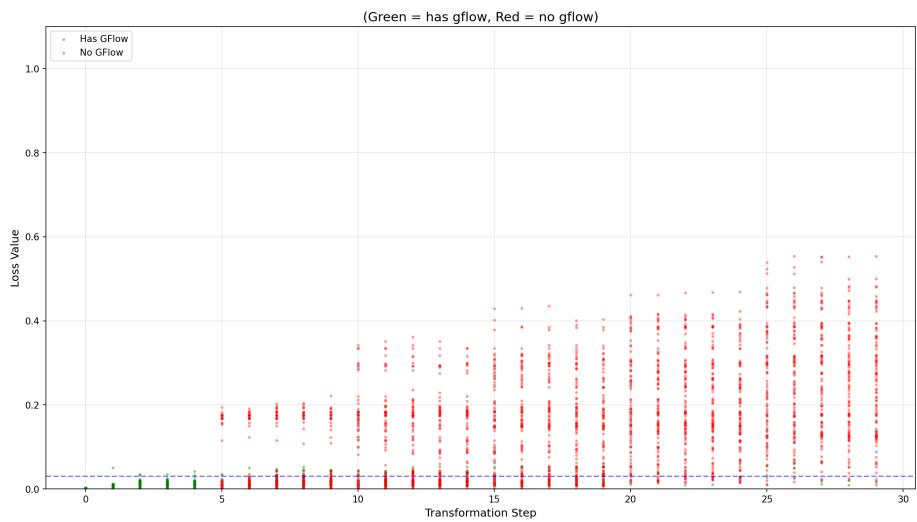


Figure 8: GFlow loss transition scatter plot showing the relationship between step number and loss, with colored gflow.